Attention Is All You Need
Vaswani et al.

Abstract: The dominant sequence transduction models are based on complex recurrent or convolutional neural n

Section 3: Model Architecture

The Transformer follows this overall architecture using stacked self-attention and point-wise, fully connected laye

Section 4: Scaled Dot-Product Attention

We compute the matrix of outputs as softmax(QK^T / sqrt(d_k))V. Multi-head attention allows the model to jointly